

Which Organizational Learning Disability Does Healthcare Suffer From Most?

The prompt makes a strong case for “I am my position,” and I do not think it is wrong — but I would argue it is a symptom rather than the root. The disability healthcare suffers from most is **the delusion of learning from experience**, and it is the one that generates most of the others.

Senge’s point is that we learn best from experience, but we never directly experience the consequences of our most important decisions. Every organization has a **learning horizon** — a span in time and space within which we can actually observe the results of what we did. When consequences fall outside that horizon, the feedback loop that would let us learn never closes. We keep acting, we keep accumulating “experience,” and none of it teaches us anything.

Healthcare’s learning horizons are unusually short relative to the timescale on which its decisions actually resolve. A hospitalist makes a discharge decision on Tuesday; the consequence surfaces as a readmission weeks later, in a different unit, attributed to a different clinician, and never travels back to the person who made the call. Pronovost notes that heart failure patients carry roughly a 40% readmission rate at 90 days — an enormous signal that is almost perfectly invisible to the individuals generating it. The same structure holds across the industry. A nephrologist adjusts a peritoneal dialysis prescription at a monthly visit, and the cyclor writes the answer into its logs every night for four weeks, where no one reads it until the following visit. A device designer never watches sterile processing struggle with their geometry. In each case the actor behaves rationally on the information they can see, and that information is systematically incomplete.

This is why I think it sits underneath the other disabilities rather than beside them. “I am my position” is what a short learning horizon feels like from the inside — if you can only observe your own slice, defining yourself by that slice is the reasonable adaptation. “Fixation on events” follows too: when you cannot see trends because the data never returns to you, discrete events are genuinely all you have. And the boiled frog is the terminal case — degradation slow enough to fall outside everyone’s horizon is degradation nobody is positioned to notice.

A pathway. The intervention is not more training or better attitudes; it is engineering the feedback loop so that consequences return to the decision-maker while they can still learn from them. Pronovost’s work on catheter-related bloodstream infections is an existence proof: the checklist gets the attention, but the mechanism that mattered was feeding unit-level infection rates back to the specific clinicians producing them, on a cadence short enough to connect action to outcome. A problem the size of breast cancer was largely eliminated by closing a learning loop, not by generating new science.

Systems engineering formalizes exactly this. Verification and validation are institutionalized feedback loops, and requirements traceability is a mechanism for carrying consequences backward to the decision that caused them. Even the context diagram is a learning-horizon tool — drawing the boundary and enumerating what crosses it is a structured admission that effects will land outside your view, and a way to name them before they surprise you.

Three concrete moves follow. First, for any recurring decision, identify where its consequence materializes and build a route back. Second, measure trends rather than incidents, so boiled-frog cases become visible. Third, attach outcomes to decisions rather than departments — blame fragments learning, traceability preserves it.

Question for the group: in your own setting, what is a decision you make routinely whose consequence you have genuinely never observed? I suspect naming those is harder than it sounds, which is rather the point.